

Stephen Jonathan Pratama

<https://www.linkedin.com/in/stephen-jonathan-pratama-260365245>

Jakarta Utara, Jakarta 14420 • stephen.j.pratama@gmail.com • + 62 85155139952

TECHNICAL SKILLS

- HTML, CSS, Javascript, NodeJS, PHP, Python, C++, Flutter, UML, Power Bi
- MySQL, Postgres, GitLab CI/CD, AWS, Computer Vision, Natural Language Processing
- Docker, Podman, Kubernetes, Flask, SkLearn, Pytorch, Huggingface, Kafka, JWT, OAuth, Linux

EDUCATION

CALVIN INSTITUTE OF TECHNOLOGY

Jakarta, Indonesia

Seventh-semester student, IT & Big Data Analytics

2022-2026

GPA: 3.73 / 4

PROFESSIONAL EXPERIENCE

PT DIGDAYA SOLUSI TEKNOLOGI

Jakarta, Indonesia

Junior Full-Stack Developer

2025-Present

- Developed and maintained full-stack web applications using modern technologies, ensuring responsive design and optimal user experience.
- Built backend APIs with Node.js/Express and integrated with PostgreSQL/MySQL databases for efficient data management.
- Implemented frontend features using Vue.js with dynamic and interactive interfaces.
- Integrated Apache Kafka for real-time data streaming and combined with WebSocket to enable real-time vehicle tracking on the platform.
- Collaborated with cross-functional teams to design, develop, and deploy new product features on time.
- Optimized application performance, reduced load times, and improved scalability through code refactoring and best practices.
- Participated in code reviews, testing, and debugging to ensure high-quality software delivery.

RELATED PROJECTS

SILA: SECURE ONLINE VIDEO CALL SYSTEM FOR WARGA BINAAN LAPAS

Jakarta, Indonesia

PT Digdaya Solusi Teknologi - Junior Software Developer

2025

- Developed SILA (Sistem Layanan Audio Visual / Wartelsuspas), a secure online communication platform designed to enable encrypted video calls between Warga Binaan Lapas (inmates) with their families through encrypted video calls and virtual meetings.
- Built the backend system using Node.js (Express) and PostgreSQL, implementing robust authentication, session management, and data encryption.
- Integrated FaceAPI.js for real-time facial verification to authenticate Warga Binaan during login, enhancing system security and preventing impersonation.

- Utilized WebSocket for real-time video call signaling and communication, enabling low-latency, reliable connections between users.
- Implemented Google Gemini API for automatic meeting transcription and summarization, providing accessible records of conversations while maintaining confidentiality.
- Used FFmpeg to convert and process recorded meeting audio for storage, playback, and AI-based summarization pipelines.

AUTOMATED CI/CD PIPELINE WITH GITLAB, DOCKER, AND AWS EC2

Jakarta, Indonesia

PT Digdaya Solusi Teknologi - Junior Software Developer

2025

- Designed and deployed a Continuous Integration and Continuous Deployment (CI/CD) pipeline on Amazon EC2 using GitLab and Docker to automate code delivery, review, and deployment for a multi-branch development workflow.
- Set up GitLab Runner on AWS EC2 to automate build, test, and deployment stages within containerized environments.
- Configured a multi-stage pipeline (notify, build, deploy) defined in .gitlab-ci.yml to automate testing, building, and deployment.
- Implemented branch-based restrictions: developers can only push to the development branch, triggering automated Docker builds and test deployments.
- Designed an email notification system that alerts supervisors upon commits to the development branch for review.
- Deployed services as Docker containers using a private Docker registry with Docker-in-Docker (DinD) setup for isolation and build efficiency.
- Applied best practices for CI/CD security, including SSH key encryption, token-based authentication, and branch protection rules.
- Delivered a reliable and auditable deployment process that enhanced collaboration, quality control, and production stability.

FINETUNE A DEEPSEEK R1 TRANSFORMER MODEL FOR TEXT SUMMARIZATION

Jakarta, Indonesia

Student - Junior AI Engineer

2024

- Developed a text summarization system by fine-tuning the DeepSeek-R1 Distill Llama-8B large language model using Python, PyTorch, and Hugging Face Transformers, enabling automatic generation of concise, high-quality summaries for long-form news articles.
- Prepared and preprocessed diverse news datasets, including text cleaning, tokenization, and handling domain-specific terminology.
- Applied transfer learning with LoRA (Low-Rank Adaptation) to efficiently fine-tune the pretrained model, reducing GPU memory usage and training time.
- Experimented with quantization and mixed-precision training (FP16) to optimize inference latency without sacrificing summarization accuracy.
- Conducted hyperparameter tuning (learning rate, batch size, sequence length) to improve performance on the summarization task.
- Evaluated model quality using ROUGE, BLEU, and human evaluation, ensuring summaries retained key context, factual accuracy, and coherence.
- Integrated the fine-tuned model into a real-time summarization pipeline capable of processing continuous news feeds.